

Project Documentation

This project has 2 solutions as follows:

1. Whole ELT Pipeline with modern data tools
2. Local solution using jupyter notebook

Solution 1:

Overview

This project demonstrates a modern **data ingestion, transformation, and alerting pipeline** built using **AWS S3, Snowflake, dbt, Docker, Airflow, and Slack**.

The goal is to ingest raw financial data (invoices, organizations), transform it into analytics-friendly models, and trigger alerts when organization balances change by more than 50% day-over-day.

Steps:

1. Data Ingestion

- a. **IAM Setup:** Created IAM user/role in AWS for secure authentication and access control.
- b. **S3 Bucket Setup:** Configured an S3 storage bucket to upload .csv files (invoices, organizations).
- c. **File Upload:** Uploaded raw .csv datasets into the bucket.

2. Data Extraction & Load (Snowflake Setup)

- a. **Platform Setup:**
 - i. Created user: **DEEL**
 - ii. Created warehouse: **deel**
 - iii. Created database: **deel_db**
 - iv. Created schema: **raw** (for staging raw data)
 - v. Created role: **deel_ownership** with full privileges (warehouse, schema, tables).
 - vi. Assigned role **deel_ownership** to user **DEEL**.

3. Data Extraction Setup:

- a. Created a storage integration between **Snowflake** and **AWS S3**.
- b. Created a **Snowflake** stage pointing to **S3** for raw file ingestion.

4. Data Extraction:

- a. Used **COPY INTO** commands to load data into:
 - i. **deel_db.raw.invoices**
 - ii. **deel_db.raw.organizations**

5. Data Transformation (dbt Setup)

- a. **DBT Setup:**
 - i. Installed **dbt-core** locally inside VS Code.
 - ii. Initialized dbt project **deel**.
 - iii. Configured Snowflake connection in **profiles.yml**.

b. Staging Models:

- i. Created **sources.yml** mapping **deel_db.raw** as dbt source.
- ii. Created **stg_invoices** and **stg_organizations** in schema **deel_db.staging**.

c. Warehouse Models (Facts & Dimensions):

i. Facts:

- 1. **fct_org_daily_balances**(grain: date × organization)
- 2. Measures: **total_amount**, **payment_amount**, **balance_amount = total_amount - payment_amount**

ii. Dimensions:

- 1. **dim_organizations** (grain: organization)
- 2. Enriched with: total invoices, total payments, total invoice amounts.

d. Testing & Documentation:

- i. Added data tests (**not_null**, **unique**) on the primary keys of the Staging tables
- ii. Added **uniqueness and not null tests** on the grain of the **fct_org_daily_balances(date/organization)** table
- iii. Added **not null and uniqueness tests** on the **dim_organizations**

6. Alerts (Airflow + Slack)

a. Docker & Airflow Setup:

- i. Installed Docker Desktop.
- ii. Built custom Airflow image via **Dockerfile**.
- iii. Created **docker-compose.yml** to run Airflow with webserver, scheduler, and workers.

b. Airflow Connections:

- i. Configured Snowflake connection (for dbt/fct_invoices).
- ii. Configured Slack connection (Webhook).

c. Slack API Setup:

- i. Created Slack workspace and new channel **#balance-change-alert**.
- ii. Created a new Slack App, generated API token & Webhook URL.
- iii. Added token to Airflow Slack connection.

d. Alert Logic (Airflow DAG):

- i. Query **deel_db.facts.fct_org_daily_balances** for latest date only.
- ii. Compute day-over-day change in **balance_amount**.
- iii. Trigger Slack alert if balance changes > 50%.
- iv. Post formatted alert message in **#balance-change-alert**.

7. Best Practices / Extras

- a. Version controlled the entire project with GitHub.
- b. Created requirements.txt to lock dependencies.
- c. Enabled Snowflake RBAC

- d. Designed dbt models using a layered architecture: RAW → STAGING → FACTS/DIMS → ALERTS.
- e. Added logging in Python alert script in Airflow.

Future Scope

- 1. All secrets (Snowflake, Slack) stored in Airflow's encrypted backend.
- 2. Add CI/CD pipeline for SQL + DAG deployment.

Solution 2:

This is a .ipynb file that performs all actions from solution 1. Instead of using modern tools, the solution is created using a jupyter notebook.

Overview:

1. Environment Setup

- Imports Python libraries: pandas, os, json, requests, datetime
- Reads environment variables:
 - INVOICES_CSV, ORGS_CSV → paths to CSVs
 - SLACK_WEBHOOK_URL → Slack webhook URL

2. Data Ingestion

- If CSV paths are set, loads invoices & organizations data from local files.
- Otherwise, throw an error.
- Saves them into a local output folder

3. Staging Layer (Mimicking dbt staging)

- Cleans column names, standardizes date formats.
- Creates **stg_invoices** and **stg_organizations** DataFrames.

4. Warehouse Layer (Facts & Dimensions)

- **Fact Table (fct_invoices)**
 - Aggregated by (organization_id, invoice_date)

- Metrics:
 - `total_amount = sum(amount)`
 - `total_payment = sum(payment_amount)`
 - `balance = total_amount - total_payment`
- **Dimension Table (dim_organizations)**
 - Aggregated by `organization_id`
 - Metrics:
 - `total_invoices`
 - `total_invoice_amount`
 - `total_payment_amount`
- Saves both tables as CSVs in the output folder.

5. Data Quality Checks

- Ensures:
 - No duplicate invoice IDs
 - No negative amounts or payments
 - Valid organization references

Prints warnings if issues are found.

6. Alerts Logic (Day-over-Day Balance Change)

- Looks at `fct_invoices` for the **latest available date**
- Computes % change in balance for each organization:

$$\text{pct_change} = (\text{latest_balance} - \text{prev_balance}) / (\text{prev_balance}) * 100$$
- If absolute % change > **50%**, trigger an alert.

7. Alert Delivery

- If `SLACK_WEBHOOK_URL` is set:
 - Sends formatted JSON message to Slack channel `#balance-change-alert`.

8. Outputs

- Writes following CSVs into the output folder:
 - stg_invoices.csv
 - stg_organizations.csv
 - fct_invoices.csv
 - dim_organizations.csv
- sends alerts on Slack